

Graph-based Neuron Agglomeration using 3D Skeletons

Anonymous CVPR submission

Paper ID 446

Abstract

Advancements in electron microscopy image acquisition have created massive connectomics datasets which are petabytes in size and make manual segmentation not feasible. Proposed methods for automatic segmentations generate super-pixel approximations of neural membranes followed by agglomeration strategies to create full neuron reconstructions. However, pixel-based segmentation results widely neglect global geometric properties. We generate skeletons of membrane labels to approximate neural pathways in 3D. This allows us to efficiently apply graph-based optimization strategies and to train automatic classifiers for shape description against manually labeled ground truth. Our classifiers detect feasible and impossible neural pathways by looking at geometric properties alone, and are able to improve the segmentation labels accordingly. We demonstrate the performance of our classifiers on multiple real-world connectomics datasets with average variation of information improvement of $X\times$. We discuss our findings and provide insights to learning shape features for neuron agglomeration.

1. Introduction

The field of connectomics is concerned with reconstructing the wiring diagram of the brain at nanometer resolutions. Recent advancements in image acquisition using multi-beam serial-section electron microscopy (sSEM) has allowed neuroscientists to produce terabytes of electron microscopy (EM) image data every hour [12]. Neuroscientists believe that reconstructing an entire mammalian brain at fine resolution will enable new insights into the workings of the brain [16]. These observations will allow for new advancements in neuromedicine and artificial intelligence (CITE). Segmentation is part of this reconstruction process, assigning a unique label for every neuron in the EM image. It is not feasible for domain experts to manually segment the vast amount of 3-D image data to model an entire brain.

A significant amount of research focuses on automatic reconstruction of the neurons in EM images because of

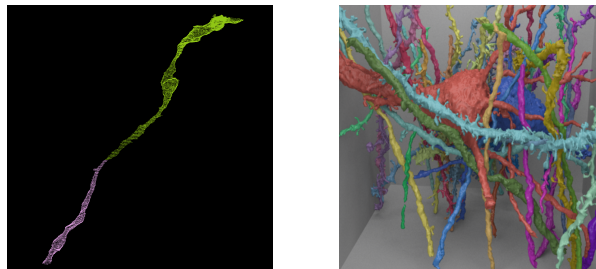


Figure 1: We use 3D skeletons to combine neuron labels to full reconstructions.

the scope and importance of the problem (ADD REFERENCES). All of these algorithms extract neuronal processes through the 3-D volumes using only the raw image data as input. Oftentimes, convolutional neural networks (CNNs) predict membrane probabilities or affinities between voxels and apply simple thresholds to agglomerate the voxels into clusters [20, 29]. These *per-voxel* algorithms produce excellent results but currently fall short of complete reconstructions with error rates of approximately $X\%$.

Researches currently address the failures of the *per-voxel* algorithms by training random-forest classifiers to agglomerate an oversegmentation of voxels [23] (CITE NEUROPROOF). These classifiers take the output of the *per-voxel* algorithms as input and generate high-level statistics such as affinity distributions between pixel regions. Presently, these methods use hand-designed features despite the evidence that machine-learned features perform better [4]. These *per-supervoxel* algorithms outperform the *per-voxel* algorithms but still do not provide the accuracy needed for large scale reconstructions of the brain. Additionally, these methods do not fully leverage the wealth of shape information available. E.g., it is well known that neurons angle at less than 30 degrees (CITE).

Here we present a *graph-based* strategy that builds on top of the outputs of *per-supervoxel* methods by taking neuronal shape properties into account. Similarly we take as input automatic segmentations. However, we focus on general neuron shapes, traversing through the label volumes to

identify potential merge candidates. From these locations we extract machine-learned features and generate a probability that two label segments should be joined. We construct a graph representation of the input label volume. Using graph-based optimization strategies enables us to enforce global constraints that more closely match the underlying biology of the images. Lastly, we apply a global segmentation algorithm to produce a better reconstruction.

2. Related Work

A large amount of connectomics research considers the problem of extracting segmentation information at the voxel level from the raw EM images. Some methods apply computationally expensive graph partitioning techniques with a node per voxel [1] following the previous work in computer vision using normalized cuts for image segmentation [15, 31, 32]. Meanwhile, others focus on training 2-D convolutional neural networks to predict membrane probabilities per image slice [6]. More recent strategies augment these neural networks with the z -dimension creating full 3-D convolutional networks [20]. Oftentimes these networks now produce probabilities for the affinity between voxels rather than probabilities for a voxel belonging to a cell membrane [29]. An extensive amount of research in computer vision, applied mathematics, and statistics evaluates the loss functions and optimizers for these networks [5, 21, 22].

These above neural networks generate probabilities that neighboring voxels belong to the same neuron. Many algorithms work at a level above these networks and train random-forest classifiers to produce segmentations of the EM images where every unique neuron in the volume has a unique label [19, 23, 27, 34] (CITE NEUROPROOF). Despite the success of these classifiers, they often make mistakes based on the local information leading to errors in the segmentation. Proofreading methods create a “human-in-the-loop” framework that allows users to find errors and correct them [9, 10, 11, 35]. More recently, flood-filling networks merge the process of generating voxel affinities and then producing segmentations by training an end-to-end neural network that outputs segmentations [14]. Although these networks produce impressive segmentation accuracies, they are currently too slow for large scale connectomics reconstructions.

Many segmentation and clustering algorithms use graph partitioning techniques [1]. The multicut graph partitioning algorithm extracts segments the graph and remove edges such that there are no cycles. This closely resembles the biological constraint that neurons have a topological genus of zero. There are several useful multicut heuristics which provide good approximations with reasonable computational costs [13, 17, 18]

A sizable amount of computer vision research analyzes

various hand-designed shape features for determining similarity between objects [24] and for segmentation[7]. Currently, the random-forest classifiers used in connectomics rely on hand-designed features to make merge decisions. Some of these features encode the distribution of probabilities that voxels along a segment boundary belong to the same neuron as their neighbors [23, 28]. However, there is evidence that machine-learned features outperform hand-designed ones [4].

The fields of computer graphics, mathematics, and biomedical visualization produce extensive research into the skeletonization of 3D binary volumes. Some of the research explores topologically consistent thinning algorithms and potential medical applications of these methods [25, 26]. In computer graphics, fast medial axis algorithms allow animators to quickly move characters through successive frames [2, 3]. In these works, the skeleton acts as the simplest representation of a 3-D volume. The Tree-structure Extraction Algorithm for Accurate and Robust Skeletons (TEASER) has been used by connectomics researchers to parameterize 3-D shapes [30, 33].

3. Method

Our method takes as input an existing oversegmentation of an *EM* image volume. Section ?? discusses two pre-existing pipelines for generating these segmentations. We evaluate our proposed method on the outputs from both of these pipelines.

3.1. Graph Creation

We need to generate nodes N and edges E to apply a graph-based optimization strategy for segmentation. In addition, these edges need non-negative weights.

3.1.1 Node Generation

The simplest node generation strategy creates one node for every unique segment label in the input volume. However, there are often a small number of labels in the volume corresponding to very small regions. It is difficult to extract useful shape features from these segments because of their small, and often random, shape. These segments are pruned from the graph and do not have corresponding nodes. For the proposed framework, we remove segments with fewer than 20,000 voxels (ADD PERCENT, current estimate less than 10%). Figure 2 shows two typical input segments that receive a corresponding node in the graph. .

3.1.2 Edge Generation

A naïve approach to generating edges simply produces an edge between all segments that have a single pair of neighboring voxels. Current agglomeration strategies such as




Figure 2: Two example outputs of the TEASER skeletonization algorithm.

NeuroProof or GALA use this criteria when deciding pairs of segments to merge into one neuron. However, for us this method will produce too many edges in our graph, many of which are easily prunable. Many of the segments that are erroneously split have a similar structure. Consider Figure 3 which shows two such pairs. In both instances the segments around the break follow the same general shape. The segment is tubular in the close vicinity with an abrupt break perpendicular to the elongated direction. We present the following algorithm to identify these locations.

We generate a skeleton for every segment following the intuition that a skeleton represents a simplified representation of the overall shape of an object. We use the Tree-structure Extraction Algorithm for Accurate and Robust Skeletons (TEASER) [30, 33]. These skeletons consist of a sequence of *joints* with edges between successive joints. We prune the joints that are within 50 voxels of each other to reduce unnecessary branching. For the purposes of this algorithm, joints that have only one connected neighbor are referred to as *endpoints*. Figure 2 shows the skeletons in white of two typical segments from the label volume. After skeletonization, we begin to identify segments for merge consideration with the following two-pass heuristic. In the first pass, we iterate over all endpoints e belonging to a segment S and create a set of segments S'_e which includes all labels that have a single voxel within t_{low} nanometers from e . This first pass often includes too many candidates that should not merge so we apply an additional pass to prune these sets. In the second pass, we consider all of the segments S'_e for every endpoint e . If a segment $S' \in S'_e$ has an endpoint within t_{high} nanometers of e , the segment S and S' are considered for merging. We store the midpoint between the two endpoints as the “center” of the potential merge. Algorithm (ADD REF) provides pseudocode for this edge generation algorithm. The results in this paper follow from $t_{low} = 240$ and $t_{high} = 600$.

```
function GENERATEEDGES(skeletons)
```

```
  for skeleton in skeletons do
```

```
    candidates = set()
```

```
    for endpoint in skeleton do
```




Figure 3: Two erroneously split segments that should merge together. Most segments that we want to merge have the same general structure.

TODO ADD CODE

end for

end for

end function

Note that the above algorithm does not enforce any segment adjacency constraints. Figure (ADD FIGURE) shows two examples that would not be considered in previous research since they are not adjacent. Finding and merging these examples is one of the benefits of retreating from per-pixel algorithms.

3.2. Edge Probabilities

The previous section outlines how to generate edges between nodes in our graph structure. Here we introduce a neural network architecture for generating probabilities that two nodes sharing an edge belong to the same neuron. The neural network takes as input only data from the input segmentation.

3.2.1 Network Architecture

The above skeletonization algorithm produces 3-D locations that require further consideration for merging. To determine which of these segments should actually merge, we train a 3-D convolutional neural network. We extract a cubic region with length 1200nm centered at these locations. These cubes will provide the local information for the neural network to predict which neighboring segments belong to the same neuron. A cubic length of 1200nm provides enough local shape context for the network without introducing an excessive amount of noise. Figure 4 shows cubes of lengths 800nm, 1200nm, and 1600nm.

We transform the extracted cubes for input into the neural network. The network takes three input channels for each voxel in the 3-D volume corresponding to the following mapping. Consider a pair of segments with labels l_1 and l_2 for every voxel v in the cube of interest. The first channel is 0.5 if $v = l_1$ and -0.5 otherwise, the second channel is 0.5 if $v = l_2$ and -0.5 otherwise, and the third channel




Figure 4: The white outlines show three different possible cubic sizes to input into the neural network. The left example (800nm) provides less local context than the middle example (1200nm). The right example (1600nm) extracts too large of a local region that produces noise as one of the segments leaves the bounding box only to reenter.

is 0.5 if $v = l_1$ or $v = l_2$ and -0.5 otherwise. Each 3-channel volume is subsequently downsampled into an array of size $(3, 22, 68, 68)$ using nearest-neighbor interpolation. Our neural network trains on these 4-D arrays to generate probabilities that l_1 and l_2 should merge.

Our network architecture has three layers of double convolutions followed by a max pooling step [5]. As described above, the input with into the network is $(22, 68, 68)$ with 3 channels. The filter size of the layers are 16, 32, and 64 with size doubling deeper into the network. The first two max pooling layers are anisotropic with pooling only in x and y to match the anisotropic nature of our EM datasets. All convolutions have kernel sizes of $(3, 3, 3)$. The output size of the final max pooling layer is $(5, 5, 5)$ with 64 channels. The output is flattened into a 1-D vector with 8000 entries which enters two fully connected layers with 512 and 1 dimensions in the output respectively. All activation functions are LeakyReLU with $\alpha = 0.001$ except for the final activation which uses a sigmoid function [8, 21]. There is a dropout of 0.2 after every pooling layer and the first dense layer. There is a dropout of 0.5 after the final dense layer. The above method uses stochastic gradient descent with Nesterov’s accelerated gradient [22]. This optimizer has an initial learning rate of 0.01, momentum of 0.9, and a decay rate of $5 * 10^{-8}$. Figure 5 provides an overview of the proposed architecture.

3.2.2 Data Augmentation

We apply some data augmentation to the generated examples to increase the size of the training datasets. We consider all rotations of 90 degrees along the xy -plane in addition to mirrors along the x and z axes. This produces an additional 16 times more training data.

3.3. Agglomeration

After constructing the above graph structure we can apply a graph-based segmentation strategy. There are many formulations of graph-based optimization strategies that provide different guarantees on their output. Neurons in the brain should be acyclic, i.e. the output shape should have a genus of zero. Current connectomics agglomeration techniques do not leverage this additional information but rather consider neighboring regions in successive order without regard to loop creation. In our graph formulation we can enforce this topological property by applying a multicut partition onto the graph which generates a forest on the nodes. There are several heuristics that solve the multicut problem. For our purposes we use the Kernighan-Lin algorithm [17].

4. Evaluation

We evaluate our proposed methods on two different connectomics datasets from two different species.

4.1. Datasets

4.1.1 Kasthuri

4.1.2 FlyEM

4.1.3 Segmentation Pipeline and Baseline

4.2. Skeleton Pruning

HOW DO WE EVALUATE THE PRUNING METHOD
HOW MANY EXAMPLES DO WE MISS

4.3. Classifier Training

HOW DO WE TRAIN THE CLASSIFIER - split the
data in half and have training and validation DATA AUG-
MENTATION HOW MANY LEARNABLE PARAME-




Figure 5: The architecture for the neural networks follows the VGG style of double convolutions followed by a max pooling operation. The number of filters doubles each layer leading to a fully connected layer and a sigmoid activation function.

TERS WHAT NETWORKS DID WE TRY HAVE A TABLE WITH PARAMETERS: ITERATIONS, momentum, batch size, learning rate

4.4. Graph-based Strategies

5. Results

5.1. Skeleton Pruning

5.1.1 Kasthuri

5.1.2 FlyEM

5.2. Classification Performance

5.2.1 Kasthuri

5.2.2 FlyEM

5.3. Graph Based Strategies

5.3.1 Kasthuri

5.3.2 FlyEM

6. Conclusions

References

- [1] B. Andres, T. Kroeger, K. L. Briggman, W. Denk, N. Korogod, G. Knott, U. Koethe, and F. A. Hamprecht. Globally optimal closed-surface segmentation for connectomics. In *European Conference on Computer Vision*, pages 778–791. Springer, 2012. 2
- [2] I. Baran and J. Popović. Automatic rigging and animation of 3d characters. In *ACM Transactions on Graphics (TOG)*, volume 26, page 72. ACM, 2007. 2
- [3] G. Bharaj, T. Thormählen, H.-P. Seidel, and C. Theobalt. Automatically rigging multi-component characters. In *Computer Graphics Forum*, volume 31, pages 755–764. Wiley Online Library, 2012. 2
- [4] J. A. Bogovic, G. B. Huang, and V. Jain. Learned versus hand-designed feature representations for 3d agglomeration. *arXiv preprint arXiv:1312.6159*, 2013. 1, 2
- [5] K. Chatfield, K. Simonyan, A. Vedaldi, and A. Zisserman. Return of the devil in the details: Delving deep into convolutional nets. *arXiv preprint arXiv:1405.3531*, 2014. 2, 4
- [6] D. Ciresan, A. Giusti, L. M. Gambardella, and J. Schmidhuber. Deep neural networks segment neuronal membranes in electron microscopy images. In *Advances in neural information processing systems*, pages 2843–2851, 2012. 2
- [7] R. W. Connors, M. M. Trivedi, and C. A. Harlow. Segmentation of a high-resolution urban scene using texture operators. *Computer Vision, Graphics, and Image Processing*, 25(3):273–310, 1984. 2
- [8] K.-I. Funahashi. On the approximate realization of continuous mappings by neural networks. *Neural networks*, 2(3):183–192, 1989. 4
- [9] D. Haehn, J. Hoffer, B. Matejek, A. Suissa-Peleg, A. K. Al-Awami, L. Kamensky, F. Gonda, E. Meng, W. Zhang, R. Schalek, et al. Scalable interactive visualization for connectomics. In *Informatics*, volume 4, page 29. Multidisciplinary Digital Publishing Institute, 2017. 2
- [10] D. Haehn, V. Kaynig, J. Tompkin, J. W. Lichtman, and H. Pfister. Guided proofreading of automatic segmentations for connectomics. *arXiv preprint arXiv:1704.00848*, 2017. 2
- [11] D. Haehn, S. Knowles-Barley, M. Roberts, J. Beyer, N. Kasthuri, J. W. Lichtman, and H. Pfister. Design and evaluation of interactive proofreading tools for connectomics. *IEEE Transactions on Visualization and Computer Graphics*, 20(12):2466–2475, 2014. 2
- [12] D. G. C. Hildebrand, M. Cicconet, R. M. Torres, W. Choi, T. M. Quan, J. Moon, A. W. Wetzel, A. S. Champion, B. J.

- Graham, O. Randlett, et al. Whole-brain serial-section electron microscopy in larval zebrafish. *Nature*, 545(7654):345–349, 2017. 1
- [13] A. Hornáková, J.-H. Lange, and B. Andres. Analysis and optimization of graph decompositions by lifted multicuts. In *International Conference on Machine Learning*, pages 1539–1548, 2017. 2
- [14] M. Januszewski, J. Maitin-Shepard, P. Li, J. Kornfeld, W. Denk, and V. Jain. Flood-filling networks. *arXiv preprint arXiv:1611.00421*, 2016. 2
- [15] J. H. Kappes, M. Speth, G. Reinelt, and C. Schnörr. Higher-order segmentation via multicuts. *Computer Vision and Image Understanding*, 143:104–119, 2016. 2
- [16] N. Kasthuri, K. J. Hayworth, D. R. Berger, R. L. Schalek, J. A. Conchello, S. Knowles-Barley, D. Lee, A. Vázquez-Reina, V. Kaynig, T. R. Jones, et al. Saturated reconstruction of a volume of neocortex. *Cell*, 162(3):648–661, 2015. 1
- [17] B. W. Kernighan and S. Lin. An efficient heuristic procedure for partitioning graphs. *The Bell system technical journal*, 49(2):291–307, 1970. 2, 4
- [18] M. Keuper, E. Levinkov, N. Bonneel, G. Lavoué, T. Brox, and B. Andres. Efficient decomposition of image and mesh graphs by lifted multicuts. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 1751–1759, 2015. 2
- [19] S. Knowles-Barley, V. Kaynig, T. R. Jones, A. Wilson, J. Morgan, D. Lee, D. Berger, N. Kasthuri, J. W. Lichtman, and H. Pfister. Rhoanet pipeline: Dense automatic neural annotation. *arXiv preprint arXiv:1611.06973*, 2016. 2
- [20] K. Lee, A. Zlateski, V. Ashwin, and H. S. Seung. Recursive training of 2d-3d convolutional networks for neuronal boundary prediction. In *Advances in Neural Information Processing Systems*, pages 3573–3581, 2015. 1, 2
- [21] A. L. Maas, A. Y. Hannun, and A. Y. Ng. Rectifier nonlinearities improve neural network acoustic models. In *Proc. ICML*, volume 30, 2013. 2, 4
- [22] Y. Nesterov. A method of solving a convex programming problem with convergence rate $o(1/k^2)$. 2, 4
- [23] J. Nunez-Iglesias, R. Kennedy, S. M. Plaza, A. Chakraborty, and W. T. Katz. Graph-based active learning of agglomeration (gala): a python library to segment 2d and 3d neuroimages. *Frontiers in neuroinformatics*, 8, 2014. 1, 2
- [24] R. Osada, T. Funkhouser, B. Chazelle, and D. Dobkin. Shape distributions. *ACM Transactions on Graphics (TOG)*, 21(4):807–832, 2002. 2
- [25] K. Palágyi. A 3d 3-subiteration thinning algorithm for medial surfaces. In *International Conference on Discrete Geometry for Computer Imagery*, pages 406–418. Springer, 2000. 2
- [26] K. Palágyi, E. Balogh, A. Kuba, C. Halmai, B. Erdőhelyi, E. Sorantin, and K. Hasegger. A sequential 3d thinning algorithm and its medical applications. In *Biennial International Conference on Information Processing in Medical Imaging*, pages 409–415. Springer, 2001. 2
- [27] T. Parag, F. Tschopp, W. Grisaitis, S. C. Turaga, X. Zhang, B. Matejek, L. Kamensky, J. W. Lichtman, and H. Pfister. Anisotropic em segmentation by 3d affinity learning and agglomeration. *arXiv preprint arXiv:1707.08935*, 2017. 2
- [28] X. Ren and J. Malik. Learning a classification model for segmentation. In *null*, page 10. IEEE, 2003. 2
- [29] O. Ronneberger, P. Fischer, and T. Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical Image Computing and Computer-Assisted Intervention*, pages 234–241. Springer, 2015. 1, 2
- [30] M. Sato, I. Bitter, M. A. Bender, A. E. Kaufman, and M. Nakajima. Teasar: Tree-structure extraction algorithm for accurate and robust skeletons. In *Computer Graphics and Applications, 2000. Proceedings. The Eighth Pacific Conference on*, pages 281–449. IEEE, 2000. 2, 3
- [31] J. Shi and J. Malik. Normalized cuts and image segmentation. *IEEE Transactions on pattern analysis and machine intelligence*, 22(8):888–905, 2000. 2
- [32] S. Tatiraju and A. Mehta. Image segmentation using k-means clustering, em and normalized cuts. *Department of EECS*, 1:1–7, 2008. 2
- [33] T. Zhao and S. M. Plaza. Automatic neuron type identification by neurite localization in the drosophila medulla. *arXiv preprint arXiv:1409.1892*, 2014. 2, 3
- [34] A. Zlateski and H. S. Seung. Image segmentation by size-dependent single linkage clustering of a watershed basin graph. *arXiv preprint arXiv:1505.00249*, 2015. 2
- [35] J. Zung, I. Tartavull, and H. S. Seung. An error detection and correction framework for connectomics. *CoRR*, abs/1708.02599, 2017. 2